



Advanced Scheduling

Dr. Naser Al Madi



Project survey results

I appreciate your feedback, all of it!

<https://docs.google.com/forms/d/1JXplSwXmkMHn8MJHQ2KVyo3v0telkOP9snNHCFteN8/edit#responses>




Learning objective

- Multilevel Feedback Queue
- Multi-CPU scheduling
- $O(1)$ Scheduler
- Completely Fair Scheduling
- Project 2

Multilevel Feedback Queue (MLFQ)

Waiting

Process	CPU,I/O,CPU	Arrival	Queue
P1	5 1 , 6 , 7	4 28	1 2
P2	4 , 2 , 3	3 10	1
P3	2 , 3 , 4	4 13	1
P4	5 1 , 2 , 7	14 25	1 2



P1	P2	P3	P4	P2	P3	P1	P4	idle	P4	P1	
0	4	8	10	14	17	21	22	23	25	32	39

Multilevel Feedback Queue (MLFQ)

Response time – amount of time from when a request was submitted until the first response is produced.

	Response time	Wait time	Turnaround time
P1			
P2			
P3			
P4			

Response time – amount of time from when a request was submitted until the **first response** is produced.

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	first - arrival		
P2			
P3			
P4			



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0		
P2			
P3			
P4			

P1	P2	P3	P4	P2	P3	P1	P4	idle	P4	P1	
0	4	8	10	14	17	21	22	23	25	32	39

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0		
P2	1		
P3			
P4			

P1	P2	P3	P4	P2	P3	P1	P4	idle	P4	P1	
0	4	8	10	14	17	21	22	23	25	32	39

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0		
P2	1		
P3	4		
P4			



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0		
P2	1		
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0		
P2	1		
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0	0 + 17 + 4	
P2	1		
P3	4		
P4	3		




(6 I/O + 4 waiting in the ready queue)

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	0 + 17 + 4	
P2	1	1 +	
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4		
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4	$4 + (7 - 3)$	
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4	$4 + (7 - 3)$	
P4	3		

P1	P2	P3	P4	P2	P3	P1	P4	idle	P4	P1	
0	4	8	10	14	17	21	22	23	25	32	39

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4	$4 + (7 - 3)$	
P4	3		



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	
P2	1	$1 + (6 - 2)$	
P3	4	$4 + (7 - 3)$	
P4	3	$3 + 8 + 0$	



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	$21 + 18 = 39$
P2	1	$1 + (6 - 2)$	$5 +$
P3	4	$4 + (7 - 3)$	$8 +$
P4	3	$3 + 8 + 0$	$11 +$

P1	P2	P3	P4	P2	P3	P1	P4	idle	P4	P1	
0	4	8	10	14	17	21	22	23	25	32	39

Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	$21 + 18 = 39$
P2	1	$1 + (6 - 2)$	$5 + 9 = 14$
P3	4	$4 + (7 - 3)$	$8 +$
P4	3	$3 + 8 + 0$	$11 +$



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1


	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	$21 + 18 = 39$
P2	1	$1 + (6 - 2)$	$5 + 9 = 14$
P3	4	$4 + (7 - 3)$	$8 + 9 = 17$
P4	3	$3 + 8 + 0$	$11 +$



Multilevel Feedback Queue (MLFQ)

Process	CPU,I/O,CPU	Arrival	Queue
P1	5, 6, 7	0	1
P2	4, 2, 3	3	1
P3	2, 3, 4	4	1
P4	5, 2, 7	7	1

	Response time	Wait time	Turnaround time
P1	0	$0 + 17 + 4$	$21 + 18 = 39$
P2	1	$1 + (6 - 2)$	$5 + 9 = 14$
P3	4	$4 + (7 - 3)$	$8 + 9 = 17$
P4	3	$3 + 8 + 0$	$11 + 14 = 25$




Multilevel Feedback Queue Evaluation

- A process can move between the various queues; aging/decay can be implemented this way
- Multilevel-feedback-queue scheduler is defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to **upgrade** a process
 - method used to determine when to **demote** a process
 - method used to determine which queue a process will enter when that process needs service


Basic Multi-CPU scheduling

Multi-CPU scheduling (idea 1)




- **Run scheduler on one CPU and schedule processes on others!**
- Limitation?

Multi-CPU scheduling (idea 1)




- **Run scheduler on one CPU and schedule processes on others!**
- Limitation: bottle neck where all processors wait for a single processor to make decisions.

Multi-CPU scheduling (idea 2: symmetrical)




- **Each processor runs its own scheduler** (symmetrical scheduling scheme)

Multi-CPU scheduling (idea 2: symmetrical)



- **Each processor runs its own scheduler** (symmetrical scheduling scheme)
- Come in two variants:
 - Global queue
 - Local queues

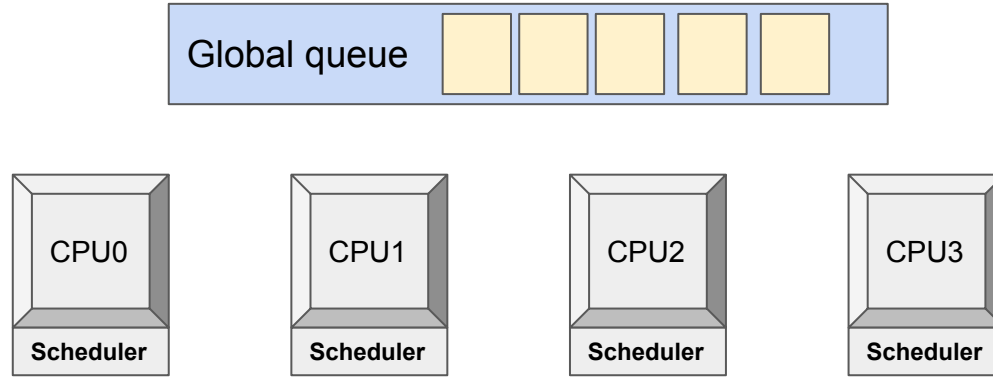
Symmetrical scheduling with global queue



Used in:
Linux 2.4
Xv6

- Advantage:
 - Good CPU utilization.
 - Fairness to all processes.
- Disadvantage:
 - The queue is the bottle neck (synched between processors).
 - Scheduler needs to be really light, since it is synched.
 - No Processor affinity (choosing exact CPU for process).

Symmetrical scheduling with global queue




Used in:
Linux 2.4
Xv6

- Advantage:
 - Good CPU utilization.
 - Fairness to all processes.
- Disadvantage:
 - The queue is the bottle neck (synched between processors).
 - Scheduler needs to be really light, since it is synched.
 - No Processor affinity (choosing exact CPU for process).

Biggest problem:


This scheme is not scalable.
Synchronization means serialization, defeating the purpose of having multiple CPUs.

Symmetrical scheduling with local queue




- Advantage:
 - Scalable (no need for synchronization).
 - Locality (TLB and registers)
- Disadvantage:
 - ?

Symmetrical scheduling with local queue



- Advantage:
 - Scalable (no need for synchronization).
 - Locality (TLB and registers)
- Disadvantage:
 - No fairness (load balancing problems) 😞


Symmetrical scheduling with local queue



- Advantage:
 - Scalable (no need for synchronization).
 - Locality (TLB and registers)
- Disadvantage:
 - No fairness 😞

The CPU on which a process runs is selected statically before it starts.

Hybrid Symmetrical scheduling




Used in:
Linux 2.6

- Advantage:
 - Global queue ensures load balancing.
 - Local queues ensure locality.

The CPU on which a process runs is selected statically before it starts.

Load balancing

- Push migration:
 - Periodically check processors loads.
 - Add work from global queue to free CPUs.
 - Redistribute all jobs among CPUs.
- Pull migration:
 - Idle processors pull tasks from busy processors.
- Migration is expensive:
 - CPU memory like TLB needs to be repopulated.
 - Migration takes time.



Modern scheduling

Evolution

Extra reading:

Understanding the Linux Kernel, Third Edition 3rd Edition by Daniel P. Bovet

Scheduling Algorithm	Operating System
O(n)	Linux 2.4 to 2.6
O(1)	Linux 2.6 to 2.6.22
CFS	Linux 2.6.23 onwards (2007 - 2016)

Process Classification in Linux

- Real time
 - Strict deadlines.
 - Should never be blocked by a low priority process.
- Normal Processes
 - Interactive
 - Constantly interact with their users.
 - When input is received, the process must wake up quickly (with small delay).
 - Batch
 - Do not require any user interaction, often run in the background.

Process Classification in Linux

- Real time
 - Strict deadlines.
 - Should never be blocked by a low priority process.
- Normal Processes
 - Interactive
 - Constantly interact with their users.
 - When input is received, the process must wake up quickly (with small delay).
 - Batch
 - Do not require any user interaction, often run in the background.

A real time process is
always a real time process

Process Classification in Linux

- Real time
 - Strict deadlines.
 - Should never be blocked by a low priority process.
- Normal Processes
 - Interactive
 - Constantly interact with their users.
 - When input is received, the process must wake up quickly (with small delay).
 - Batch
 - Do not require any user interaction, often run in the background.

A normal process could act as an interactive process and a batch process interchangeably.

Linux determines type based on **heuristics**

$O(n)$ Scheduler

- At every context switch
 - Scan the list of runnable processes
 - Compute priorities
 - Select the best process to run
- $O(n)$ performance = not scalable
 - Scalability issue observed with JVM creating many tasks
- Used a global run-queue
 - Synchronization needed (again not scalable)

Queue of Ready Processes



O(1) Scheduler

- Constant time required to pick the next process to execute.
 - Easily scales to large number of processes.
- Processes divided into 2 types:
 - **Real time:** priorities from 0 to 99.
 - **Normal processes:** priorities from 100 to 139
 - interactive
 - batch

O(1) Scheduler

- Uses dynamic priority (heuristic)
- Uses variable quantum for each process (heuristic)
- All operations done in constant time
- Combines priority, multilevel feedback queues, round robin, and heuristics to determine priority and quantum.



O(1) Scheduler

- Constant time required to pick the next process to execute.
 - Easily scales to large number of processes.
- Processes divided into 2 types:
 - **Real time:** priorities from 0 to 99.
 - **Normal processes:** priorities from 100 to 139
 - interactive
 - batch

Priority	
low	139
	...
	100
	99
High	...
	0


Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.



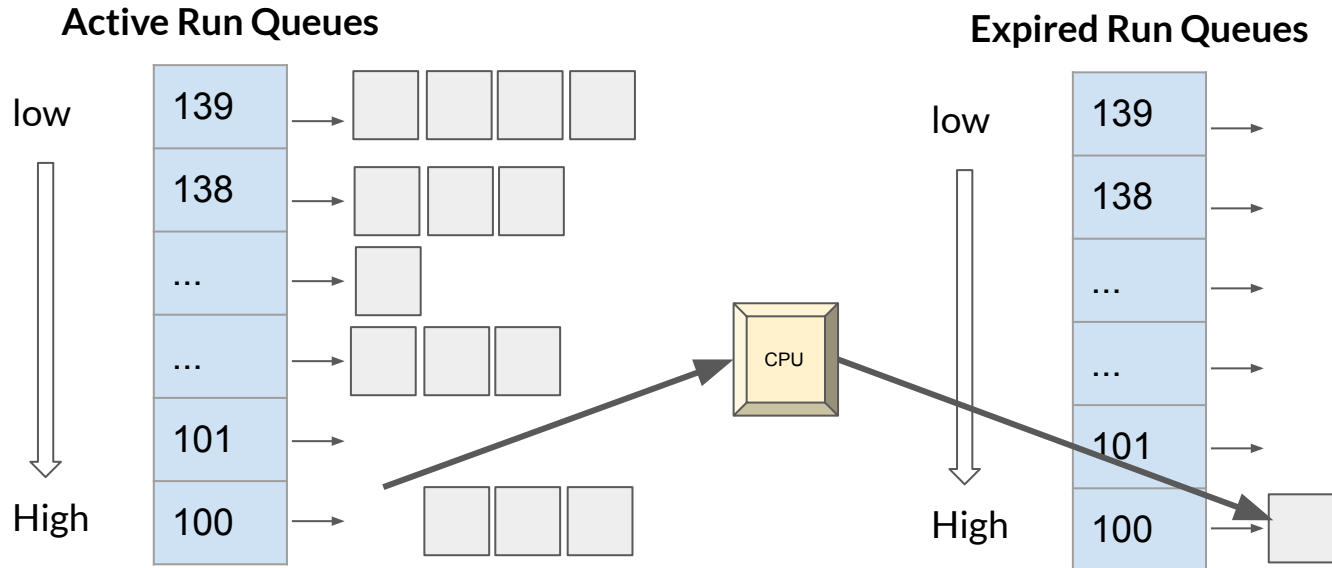
Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.



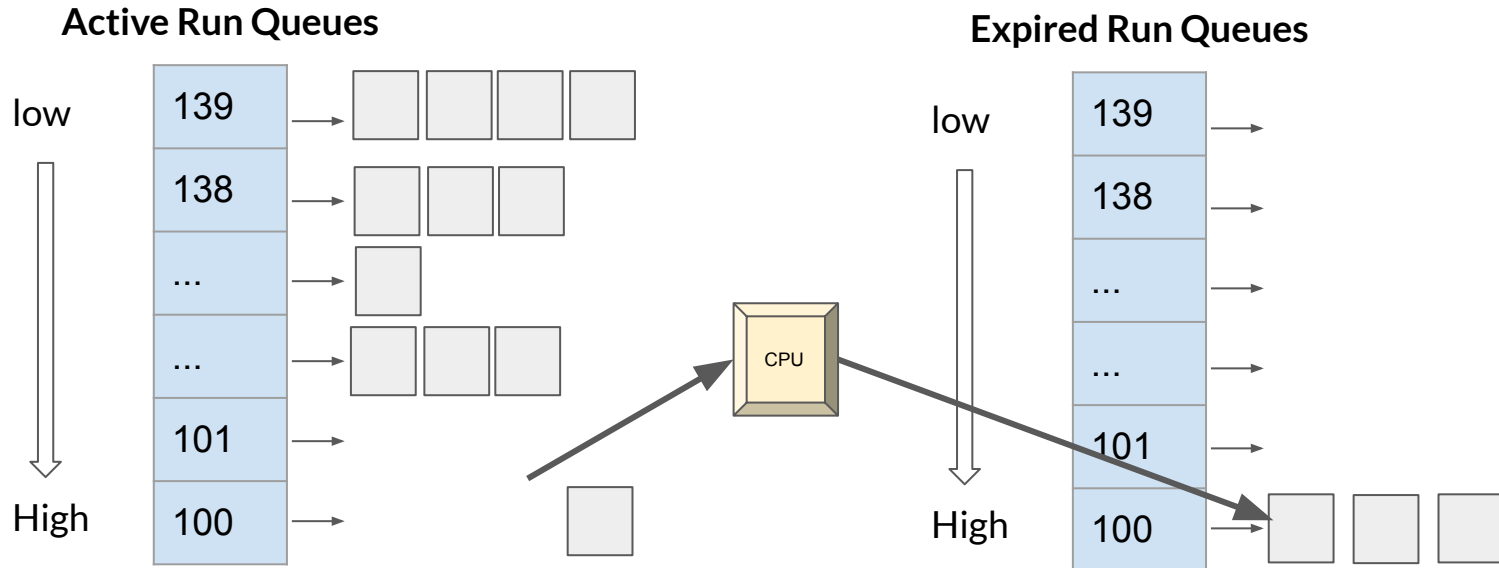
Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.



Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




This continues for a while....




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.




Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.



Scheduling Normal Processes

- Two ready queues for each CPU
 - Each queue has 40 priority classes (100 - 139)
 - 100 highest priority, 139 is lowest.



More on priorities

- Static priorities
 - 120 is the default priority
 - Nice: command line to change default priority of a process.
Nice -n [NUMBER] ./a.out

NUMBER is a value from +19 (nice) to -20 (selfish)

- Dynamic priorities (heuristic)

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(\text{static_priority} - \text{bonus} + 5, 139))$$

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(\text{static_priority} - \text{bonus} + 5, 139))$$

Bonus has value between 0 and 10

If **bonus** < 5, implies less interaction with the user, thus CPU bound process. Priority is decreased in this case.

If **bonus** > 5, implies more interaction with the user, thus more of an interactive process. Priority is increased in this case.

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(\text{static_priority} - \text{bonus} + 5, 139))$$

e.g.:

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(120 - 8 + 5, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(115, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, 115)$$

$$\text{Dynamic_priority} = 115$$

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

Dynamic_priority = MAX(100, MIN(static_priority - bonus + 5, 139))

e.g.:

Dynamic_priority = MAX(100, MIN(120 - 8 + 5, 139))

Dynamic_priority = MAX(100, MIN(115, 139))

Dynamic_priority = MAX(100, 115)

Dynamic_priority = 115 < **120**

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(\text{static_priority} - \text{bonus} + 5, 139))$$

e.g.:

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(120 - 1 + 5, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(124, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, 124)$$

$$\text{Dynamic_priority} = 124$$

Dynamic Priority

- To distinguish between batch and interactive processes
- Uses a “bonus” which changes based on a heuristic

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(\text{static_priority} - \text{bonus} + 5, 139))$$

e.g.:

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(120 - 1 + 5, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, \text{MIN}(124, 139))$$

$$\text{Dynamic_priority} = \text{MAX}(100, 124)$$

$$\text{Dynamic_priority} = 124 > \mathbf{120}$$


How does the OS set the bonus?

- Based on average sleep time:
 - I/O bound process will sleep more therefore should get a higher priority.
 - CPU bound process will sleep less, therefore should get lower priority.

Average sleep time	Bonus
>= 0 but < 100 ms	0
>= 100 ms but < 200 ms	1
>= 200 ms but < 300 ms	2
>= 300 ms but < 400 ms	3
>= 400 ms but < 500 ms	4
>= 500 ms but < 600 ms	5
>= 600 ms but < 700 ms	6
>= 700 ms but < 800 ms	7
>= 800 ms but < 900 ms	8
>= 900 ms but < 1000 ms	9
1 second	10

Dynamic Priority and Run Queues

- Dynamic priority used to determine on which queue a process is placed.
- No starvation, and everyone waits.



Setting timeslice (quantum) for each process

- Interactive processes have high priorities.
 - But likely to not complete their timeslice.
 - Given largest timeslice to ensure that CPU burst is completed without preemption.

If priority < 120:

Timeslice = (140 - priority) * 20 milliseconds

Else:

Timeslice = (140 - priority) * 5 milliseconds

Setting timeslice (quantum) for each process

- Interactive processes have high priorities.
 - But likely to not complete their timeslice.
 - Given largest timeslice to ensure that CPU burst is completed without preemption.

If priority < 120:

Timeslice = (140 - priority) * 20 milliseconds

Else:

Timeslice = (140 - priority) * 5 milliseconds

Priority:	Static Pri	Niceness	Quantum
Highest	100	-20	800 ms
High	110	-10	600 ms
Normal	120	0	100 ms
Low	130	10	50 ms
Lowest	139	19	5 ms

Limitations of $O(1)$ Scheduler

- Heuristics too complex to distinguish process type.
- Dependence between time slice and priority.
- Priority and time slice values not uniform (fairness?).

How is it constant time?


$O(1)$ scheduler has two operations:

1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.

How is it constant time?

$O(1)$ scheduler has two operations:


1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.



How is it constant time?

$O(1)$ scheduler has two operations:


1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.



How is it constant time?

$O(1)$ scheduler has two operations:


1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.



How is it constant time?

$O(1)$ scheduler has two operations:

1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.




How is it constant time?

$O(1)$ scheduler has two operations:

1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.

How is that constant time?



How is it constant time?

$O(1)$ scheduler has two operations:


1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.

Maintain a bitmap of run queues with non-zero length

Bitmap



Active Run Queues



How is it constant time?

$O(1)$ scheduler has two operations:

1. Find the queue with the lowest number with at least one task.
2. Pick a process from the front of that queue.

Use special intel instruction ***bsfl***
Find-first-bit-set


It returns the **index** of the first non-zero bit in the bitmap, in constant time.

That index is the run queue number!

Bitmap



Active Run Queues



Go over project 2



References

- Introduction to Operating Systems (Prof. Chester Rebeiro, IIT Madras)
- Understanding the Linux Kernel, 3rd Edition By Daniel P. Bovet, Marco Cesati
- Linux Kernel Development By Robert Love